

Brandon Montalvo

Lab 3: What does rust actually do

Out of bounds read and write:

hardcoded.c:

```
#include <stdio.h>

int main() {
    int numbers[5] = {1,2,3,4,5};
    printf("%d\n", numbers[10]);

    return 0;
}
```

c_hardcoded.ll:

```
; Function Attrs: noinline nounwind optnone uwtable
define dso_local i32 @main() #0 {
entry:
    %retval = alloca i32, align 4
    %numbers = alloca [5 x i32], align 16
    store i32 0, ptr %retval, align 4
    call void @llvm.memcpy.p0.p0.i64(ptr align 16 %numbers, ptr align 16 @_const.main.numbers, i64 20, i1 false)
    %arrayidx = getelementptr inbounds [5 x i32], ptr %numbers, i64 0, i64 10
    %0 = load i32, ptr %arrayidx, align 8
    %call = call i32 (ptr, ...) @printf(ptr noundef @_str, i32 noundef %0)
    ret i32 0
}
```

dynamic.c:

```
#include <stdio.h>

int main() {
    int arr[5] = {10,11,12,13,14};
    int i;
    printf("(C) Enter index: ");
    scanf("%d", &i);
    printf("Element: %d\n", arr[i]);
    return 0;
}
```

c_dynamic.ll

```

; Function Attrs: noinline nounwind optnone uwtable
define dso_local i32 @main() #0 {
entry:
    %retval = alloca i32, align 4
    %arr = alloca [5 x i32], align 16
    %i = alloca i32, align 4
    store i32 0, ptr %retval, align 4
    call void @llvm.memcpy.p0.p0.i64(ptr align 16 %arr, ptr align 16 @_const.main.arr, i64 20, i1 false)
    %call = call i32 (ptr, ...) @printf(ptr noundef @_str)
    %call1 = call i32 (ptr, ...) @_isoc99_scanf(ptr noundef @_str.1, ptr noundef %i)
    %0 = load i32, ptr %i, align 4
    %idxprom = sext i32 %0 to i64
    %arrayidx = getelementptr inbounds [5 x i32], ptr %arr, i64 0, i64 %idxprom
    %1 = load i32, ptr %arrayidx, align 4
    %call2 = call i32 (ptr, ...) @printf(ptr noundef @_str.2, i32 noundef %1)
    ret i32 0
}

```

In the llvm for c in both hardcoded and dynamic, there is no bounds check. In the hardcoded c file, it just does a direct getelementptr to index 10. When we compile, there is a warning about the index being past the end of the array. In the dynamic C file, there is no warning during compile time and in the llvm, it loads the user input and again, accesses the array with no bounds check.

hardcoded.rs:

```

fn main() {
    let arr = [1,2,3,4,5];
    println!("{}", arr[10]);
}

```

Trying to compile hardcoded rust:

```

root@crn7571-bmonta02:/CSC429/Lab3Rust# rustc hardcoded.rs -o rust_hardcoded
error: this operation will panic at runtime
  → hardcoded.rs:3:17
  |
3 |     println!("{}", arr[10]);           ^~~~~~ index out of bounds: the length is 5 but the index is 10
  |
  = note: `#[deny(unconditional_panic)]` on by default

error: aborting due to 1 previous error

```

```

root@crn7571-bmonta02:/CSC429/Lab3Rust# rustc --emit=llvm-ir hardcoded.rs
error: this operation will panic at runtime
  → hardcoded.rs:3:17
  |
3 |     println!("{}", arr[10]);           ^~~~~~ index out of bounds: the length is 5 but the index is 10
  |
  = note: `#[deny(unconditional_panic)]` on by default

error: aborting due to 1 previous error

```

When I try to compile hardcoded rust file, it is an immediate compile error with index out of bounds. No LLVM-IR could be generated either. So, this would lead to zero runtime overhead since it is not allowed to compile in the first place

Dynamic.rs:

```
use std::io;

fn main() {
    let arr = [10, 20, 30, 40, 50];

    println!("(RUST) Enter index: ");
    let mut input = String::new();
    io::stdin().read_line(&mut input).unwrap();
    let idx: usize = input.trim().parse().unwrap();

    println!("arr[{}]= {}", idx, arr[idx]);
}
```

Rust_dynamic.ll:

The file is compiled successfully. In the LLVM-IR, bounds checks are inserted at runtime.

```
; <F as core::str::pattern::MultiCharEq>::matches
; Function Attrs: inlinehint nonlazybind uwtable
-- 
store ptr %idx, ptr %args dbg.spill, align 8, !dbg !3645
%49 = getelementptr inbounds i8, ptr %args dbg.spill, i64 8, !dbg !3645
store ptr %_20, ptr %49, align 8, !dbg !3645
#dbg_declare(ptr %args dbg.spill, !3562, !DIExpression(), !3646)
; invoke core::fmt::irt::Argument::new_display
  invoke void @_ZN4core3fmt2rt8Argument11new_display17hca9c2967ca7abe86E(ptr sret([16 x i8]) align 8 %_23, ptr align 8 %idx)
    to label %bb12 unwind label %cleanup, !dbg !3646

panic:                                     ; preds = %bb10
  %50 = load i64, ptr %idx, align 8, !dbg !3644
; invoke core::panicking::panic_bounds_check
  invoke void @_ZN4core9panicking18panic_bounds_check17hbc09f5d79f1a5789E(i64 %50, i64 5, ptr align 8 @alloc_3d0e4a3a73847fef13a56cc677ef6ac0) #15
    to label %unreachable unwind label %cleanup, !dbg !3644

unreachable;                                ; preds = %panic
unreachable
--
```

In the screenshot above, it shows “; preds = %bb10” which means that the panic label comes from bb10

```

bb10: ; preds = %"_ZN4core6result19Result$LT$T$C$E$GT$6unwr
p17h7f7e30f0d4612b05E.exit"
    store i64 %t.i4, ptr %idx, align 8, !dbg !3620
    %47 = load i64, ptr %idx, align 8, !dbg !3644
    %_21 = icmp ult i64 %47, 5, !dbg !3644
    br i1 %_21, label %bb11, label %panic, !dbg !3644

bb11: ; preds = %bb10
    %48 = load i64, ptr %idx, align 8, !dbg !3645
    %_20 = getelementptr inbounds nuw i32, ptr %arr, i64 %48, !dbg !3645
    store ptr %idx, ptr %args.dbg.spill, align 8, !dbg !3645
    %49 = getelementptr inbounds i8, ptr %args.dbg.spill, i64 8, !dbg !3645
    store ptr %_20, ptr %49, align 8, !dbg !3645
    #dbg_declare(ptr %args.dbg.spill, !3562, !DIExpression(), !3646)
; invoke core::fmt::rt::Argument::new_display
    invoke void @_ZN4core3Fmt2rt8Argument11new_display17hca9c2967ca7abe86E(ptr sret([16 x i8]) align 8 %_
3, ptr align 8 %idx)
        to label %bb12 unwind label %cleanup, !dbg !3646
--
```

Looking at bb10, there is a comparison where it checks if idx < 5 then there is a branch to either bb11 if true or else go to panic. This is the runtime overhead that rust adds for safety.

If we look at the release mode:

```

root@crn7571-bmonta02:/CSC429/Lab3Rust# rustc -O dynamic.rs -o rust_dynamic_release
root@crn7571-bmonta02:/CSC429/Lab3Rust# rustc -O --emit=llvm-ir dynamic.rs -o rust_dynamic_release.ll
root@crn7571-bmonta02:/CSC429/Lab3Rust# grep -c "panic_bounds_check" rust_dynamic_release.ll
4
```

This screenshot shows that the release mode still has bound checks. If we count how many times “panic_bounds_check” appears, it is 4 but in the debug mode, the count is 8 so the release may optimize some things differently

Integer Overflows:

C int overflow:

```

root@crn7571-bmonta02:/CSC429/Lab3Rust# ./c_int_overflow
Enter i8 values: 127
Input: 127, u8: 127, i8: 127
Enter i8 values: 1
Input: 1, u8: 128, i8: -128
Enter i8 values: 1
Input: 1, u8: 129, i8: -127
Enter i8 values: 1
Input: 1, u8: 130, i8: -126
```

```
--  
%0 = load i8, ptr %input, align 1  
%conv = sext i8 %0 to i32  
%1 = load i8, ptr %u8_val, align 1  
%conv2 = zext i8 %1 to i32  
%add = add nsw i32 %conv2, %conv  
%conv3 = trunc i32 %add to i8  
store i8 %conv3, ptr %u8_val, align 1  
%2 = load i8, ptr %input, align 1  
%conv4 = sext i8 %2 to i32
```

C just has an add instruction with no overflow detection. The result is truncated and stored directly, so it has silent wrapping on overflow.

Rust int overflow (debug):

```
root@crn7571-bmonta02:~/CSC429/Lab3Rust# ./rust_int_overflow_debug  
Enter i8 value:  
127  
Input: 127, u8: 127, i8: 127  
Enter i8 value:  
1  
  
thread 'main' (81894) panicked at int_overflow.rs:14:9:  
attempt to add with overflow  
note: run with `RUST_BACKTRACE=1` environment variable to display a backtrace
```

```
--  
%_18.0 = add i8 %47, %_17  
%_18.1 = icmp ult i8 %_18.0, %47  
br i1 %_18.1, label %panic, label %bb12  
  
bb12:                                     ; preds = %bb11  
    store i8 %_18.0, ptr %u8_val, align 1  
    %48 = load i8, ptr %i8_val, align 1  
    %49 = load i8, ptr %num, align 1  
--
```

Rust debug, it inserts an overflow check using icmp and a conditional branch to panic if overflow is detected. So, for the overhead there is an extra check and a conditional branch per addition.

Rust int overflow (release):

```
root@crn7571-bmonta02:~/CSC429/Lab3Rust# ./rust_int_overflow_release
Enter i8 value:
127
Input: 127, u8: 127, i8: 127
Enter i8 value:
1
Input: 1, u8: 128, i8: -128
Enter i8 value:
1
Input: 1, u8: 129, i8: -127
Enter i8 value:
1
Input: 1, u8: 130, i8: -126
```

```
store i8 %_0.sroa.12.0.i19, ptr %num, align 1
%90 = load i8, ptr %u8_val, align 1, !noundef !4
%91 = add i8 %90, %_0.sroa.12.0.i19
store i8 %91, ptr %u8_val, align 1
%92 = load i8, ptr %i8_val, align 1, !noundef !4
%93 = add i8 %92, %_0.sroa.12.0.i19
store i8 %93, ptr %i8_val, align 1
call void @llvm.lifetime.start.p0(i64 48, ptr nonnull %_17)
call void @llvm.lifetime.start.p0(i64 48, ptr nonnull %args)
```

In the llvm, there is no zero overflow checking it just has an add instruction. This means that it is faster but unsafe. It allows wrapping behavior without panicking.

Shadowing:

```
fn main() {
    let x = 5;
    println!("x = {}", x);

    // shadow x
    let x = 10;
    println!("x = {}", x);

    {
        let x = 20;
        println!("x (inner) = {}", x);
    }

    println!("x = {}", x);
}
```

In my code, I define x, then shadow it, and then shadow it again within a scope while printing at each step.

Using grep to find all the alloca's in the LLVM file, I found this part:

```
%x3 = alloca [4 x i8], align 4
%_13 = alloca [16 x i8], align 8
%args2 = alloca [16 x i8], align 8
%_11 = alloca [48 x i8], align 8
%x1 = alloca [4 x i8], align 4
%_5 = alloca [16 x i8], align 8
%args = alloca [16 x i8], align 8
%_3 = alloca [48 x i8], align 8
%x = alloca [4 x i8], align 4
```

Rust is creating three separate stack variables (%x, %x1, and %x3) even though I used the same name “x” in my code. So, when I exit the inner scope, %x3 goes out of scope and %x1 is still there with value 10. I think what is happening is that shadowing creates distinct stack allocations for each shadowed variable. So, each time we use ‘let x’, it creates a unique LLVM variable name. When the scope ends, the shadowed variables are dropped, but the outer shadowed variables are still there.

Summary:

Which programs compiled?

- All C programs compiled (hardcoded and dynamic) successfully.
- Rust hardcoded failed to compile with a bounds check error.
- Rust dynamic compiled successfully as well.

What errors did the rust compiler produce when it didn’t build?

- Rust error said that the operation will panic at runtime because of the index out of bounds.

Which programs allowed hardcoded OOB read/write?

- Only C allowed hardcoded out-of-bounds access.

Which programs allowed user-inputted OOB read/write?

- C allowed it with no runtime check while Rust debug and release panicked on out-of-bounds input.

Runtime Behavior in Integer Overflows:

- C: $127 + 1 = -128$

- Rust Debug: Panics
- Rust Release: $127 + 1 = -128$